

(c) Copyright A. Gerbessiotis. All rights reserved.

1 Programming Project Logistics

Warning: BEFORE YOU START WORKING WITH THE ALGORITHMIC PART OF THE PROJECT, WE URGE YOU TO START WORKING ON SOME AUXILIARY COMPONENTS, ESPECIALLY IF YOU ARE NOT FAMILIAR WITH THEM, SUCH AS COMMAND LINE PROCESSING and FILE-BASED BINARY INPUT/OUTPUT. IMPLEMENTATION ERRORS AT SUCH LEVEL CAUSE SEVERAL SUBMISSIONS TO RECEIVE 0 POINTS BECAUSE THEY MAKE ALTERNATE TESTING BY THE GRADER DIFFICULT.

STEP-1. Read carefully Handout 2 and follow all the requirements specified there.

STEP-2. When the submission archive is ready, upload it to moodle

BEFORE NOON-time of Tue 4th December, 2018

For penalties check Handout 1 (Syllabus). If you submit the day before the deadline, you would have sufficient time to fix other problems. If you want to get 0 pts ignore Handout 2.

You may do one or the other option or both. You may utilize the same language or not.

OPTION 1. Do the programming part related to Huffman coding in Java, C, or C++.

OPTION 2. Do the programming part related to Kleinberg's HITS, and Google's PageRank algorithms in Java, C, or C++.

2 OPTION 1: Huffman Coding (134 points)

Huffman coding. Implement Huffman coding as explained in class and described in the Subject notes. This assumes that you will also provide an implementation of a BinaryHeap (aka array with the required Heap properties) with operations as described in class. Non-compliance to this will automatically gain you a zero. (See Handout 2.) Your implementation must work with arbitrary files binary or not. We will test it minimally on a text file, PDF, and JPG or MPEG. Avoid getting distracted with the Wikipedia code: it is not PrP compliant, and its use will cause a violation of the NJIT Honors code (see the last page of Handout 1/Syllabus under Collaboration).

Argument filename in the command-line below is an arbitrary file-name. It can be p1610s18.pdf, or p1610s18.tex, or image1.jpg, or say 01aug25_flare.mpg, or something else. The suffix (sometimes called the extension) is what follows the dot (eg pdf, tex, jpg, mpg). Although we do not ask for your implementation to be optimal and fast we require that it is relatively efficient. Thus Huffman coding (encoding or decoding) of a 2MiB file should not take more than 15-30 seconds in Java, C, or C++ on an AFS machine with reasonable load. Moreover your implementation should result in some compression savings for reasonable test files (though do not expect any in small text files or jpg files).

As stated under RULE 1, resolve issues with command-line processing or file-based input/output early in the semester. Never prompt for input as testing would be automated to some degree. So command line processing to the specifications is a MUST. My suggestion is to always read files in binary. If you don't know what i mean, pick a pdf file build a histogram of its bytes add up its frequencies and make sure that sum is equal to the size of the file as reported by an `ls -l` in a UNIX/Linux/OSX system. Write code to copy the file into another file and compare them using `md5sum` (a Unix/Linux/etc command).

```
// Encode : henc encodes filename into filename.huf and filename gets erased
//           thus x.pdf generates x.pdf.huf and then x.pdf gets erased
// Decode : hdec decodes filename.huf into filename and filename.huf gets erased
//           if filename already exists, it gets overwritten
//           thus x.pdf.huf generates x.pdf and then x.pdf.huf gets erased
// Note : Per Handout 2 you should use hencWXYZ or hdec_WXYZ for henc and hdec
% java henc filename.pdf
% java hdec filename.pdf.huf
% ./henc filename.pdf
% ./hdec filename.pdf.huf
// EXAMPLE testing steps
% cp p1610s18.pdf file1 // copy into file1
% java henc file1 // Huffman encode file1 into file1.huf
% rm -rf file1 // If not already removed, we remove it to continue testing
% java hdec file1.huf // Huffman decode file1.huf
% diff p1610s18.pdf file1 // If different something went wrong! diff might not work
% md5sum p1610s18.pdf file1 // Should have same signatures; this is better
// EXAMPLE C or C++
% cp p1610s18.pdf file1
% ./henc file1
% rm -rf file1
% ./hdec file1.huf
% diff p1610s18.pdf file1
% md5sum p1610s18.pdf file1
```

Deliverables. Read and follow Handout 2. At a minimum for a Java implementation, `hencWXYZ.java`, `hdecWXYZ.java`, `prp_WXYZ.txt`, inside an archive `prp_WXYZ.zip` or `prp_WXYZ.tar`. Or say for a C++ implementation, `hencWXYZ.cpp`, `hdecWXYZ.cpp`, `prp_WXYZ.txt`, inside an archive `prp_WXYZ.zip` or `prp_WXYZ.tar`.

3 OPTION 2: HITS and PageRank implementations (134 points)

Implement Kleinberg's HITS Algorithm, and Google's PageRank algorithm in Java, C, or C++ as explained.

(A) Implement the HITS algorithm as explained in class/Subject notes adhering to the guidelines of Handout 2. Pay attention to the sequence of update operations and the scaling. For an example of the Subject notes, you have output for various initialization vectors. You need to implement class or function `hits` (e.g. `hitsWXYZ`). We expect that no two students will have the same last 4 digits, and no student will use `WXYZ` (as in characters W, X, Y, Z) either! For an explanation of the arguments see the discussion on PageRank to follow.

```
% java hits iterations initialvalue filename
% ./hits iterations initialvalue filename
```

(B) Implement Google's PageRank algorithm as explained below adhering also to the guidelines of Handout 2. The input for this (and the previous) problem would be a file containing a graph represented through an adjacency list representation. The command-line interface is as follows. First we have the class/binary file (eg `pgrk`). Next we have an argument that denotes the number of iterations if it is a positive integer or an errorrate for a negative or zero integer value. The next argument `initialvalue` indicates the common initial values of the vector(s) used. The final argument is a string indicating the filename that stores the input graph.

```
% ./pgrk iterations initialvalue filename // in fact pgrkWXYZ
% java pgrk iterations initialvalue filename // in fact pgrkWXYZ
```

The two algorithms are iterative. In particular, at iteration t all pagerank values are computed using results from iteration $t - 1$. The `initialvalue` helps us to set-up the initial values of iteration 0 as needed. Moreover, in PageRank, parameter d would be set to 0.85. The PageRank of vertex A depends on the PageRanks of vertices T_1, \dots, T_m incident to A , i.e. pointing to A . The contribution of T_i to the PageRank of A would be the PageRank of T_i i.e. $PR(T_i)$ divided by $C(T_i)$, where $C(T_i)$ is the out-degree of vertex T_i .

$$PR(A) = (1 - d)/n + d (PR(T_1)/C(T_1) + \dots + PR(T_m)/C(T_m))$$

The pageranks at iteration t use the pageranks of iteration $t - 1$ (synchronous update). Thus $PR(A)$ on the left is for iteration t , but all $PR(T_i)$ values are from the previous iteration $t - 1$. (In an asynchronous update, we have only one vector!) Be careful and synchronize!

In order to run the 'algorithm' we either run it for a fixed number of iterations and `iterations` determines that, or for a fixed `errorrate` (an alias for `iterations`); an `iterations` equal to 0 corresponds to a default `errorrate` of 10^{-5} . A -1, -2, etc , -6 for `iterations` becomes an `errorrate` of $10^{-1}, 10^{-2}, \dots, 10^{-6}$ respectively. At iteration t when all authority/hub/PageRank values have been computed (and auth/hub values scaled) we compare for every vertex the current and the previous iteration values. If the difference is less than **errorrate** for EVERY VERTEX, then and only then can we stop at iteration t .

Argument `initialvalue` sets the initial vector values. If it is 0 they are initialized to 0, if it is 1 they are initialized to 1. If it is -1 they are initialized to $1/N$, where N is the number of web-pages (vertices of the graph). If it is -2 they are initialized to $1/\sqrt{N}$. filename first.)

Argument `filename` describes the input (directed) graph and it has the following form. The first line contains two numbers: the number of vertices followed by the number of edges which is also the number of

remaining lines. All vertices are labeled $0, \dots, N-1$. Expect N to be less than 1,000,000. In each line an edge (i, j) is represented by $i \ j$. Thus our graph has (directed) edges $(0, 2), (0, 3), (1, 0), (2, 1)$. Vector values are printed to 7 decimal digits. If the graph has N GREATER than 10, then the values for iterations, initialvalue are automatically set to 0 and -1 respectively. In such a case the hub/authority/pageranks at the stopping iteration (i.e t) are ONLY shown, one per line. The graph below will be referred to as `samplegraph.txt`

```
4 4
0 2
0 3
1 0
2 1
```

The following invocations relate to `samplegraph.txt`, with a fixed number of iterations and the fixed error rate that determines how many iterations will run. Your code should compute for this graph the same rank values (intermediate and final). A sample of the output for the case of $N > 10$ is shown (output truncated to first 4 lines of it).

```
% ./pgrk 15 -1 samplegraph.txt
Base : 0 :P[ 0]=0.2500000 P[ 1]=0.2500000 P[ 2]=0.2500000 P[ 3]=0.2500000
Iter : 1 :P[ 0]=0.2500000 P[ 1]=0.2500000 P[ 2]=0.1437500 P[ 3]=0.1437500
Iter : 2 :P[ 0]=0.2500000 P[ 1]=0.1596875 P[ 2]=0.1437500 P[ 3]=0.1437500
Iter : 3 :P[ 0]=0.1732344 P[ 1]=0.1596875 P[ 2]=0.1437500 P[ 3]=0.1437500
Iter : 4 :P[ 0]=0.1732344 P[ 1]=0.1596875 P[ 2]=0.1111246 P[ 3]=0.1111246
Iter : 5 :P[ 0]=0.1732344 P[ 1]=0.1319559 P[ 2]=0.1111246 P[ 3]=0.1111246
Iter : 6 :P[ 0]=0.1496625 P[ 1]=0.1319559 P[ 2]=0.1111246 P[ 3]=0.1111246
Iter : 7 :P[ 0]=0.1496625 P[ 1]=0.1319559 P[ 2]=0.1011066 P[ 3]=0.1011066
Iter : 8 :P[ 0]=0.1496625 P[ 1]=0.1234406 P[ 2]=0.1011066 P[ 3]=0.1011066
Iter : 9 :P[ 0]=0.1424245 P[ 1]=0.1234406 P[ 2]=0.1011066 P[ 3]=0.1011066
Iter : 10 :P[ 0]=0.1424245 P[ 1]=0.1234406 P[ 2]=0.0980304 P[ 3]=0.0980304
Iter : 11 :P[ 0]=0.1424245 P[ 1]=0.1208259 P[ 2]=0.0980304 P[ 3]=0.0980304
Iter : 12 :P[ 0]=0.1402020 P[ 1]=0.1208259 P[ 2]=0.0980304 P[ 3]=0.0980304
Iter : 13 :P[ 0]=0.1402020 P[ 1]=0.1208259 P[ 2]=0.0970858 P[ 3]=0.0970858
Iter : 14 :P[ 0]=0.1402020 P[ 1]=0.1200230 P[ 2]=0.0970858 P[ 3]=0.0970858
Iter : 15 :P[ 0]=0.1395195 P[ 1]=0.1200230 P[ 2]=0.0970858 P[ 3]=0.0970858
```

```
% ./pgrk -3 -1 samplegraph.txt
Base : 0 :P[ 0]=0.2500000 P[ 1]=0.2500000 P[ 2]=0.2500000 P[ 3]=0.2500000
Iter : 1 :P[ 0]=0.2500000 P[ 1]=0.2500000 P[ 2]=0.1437500 P[ 3]=0.1437500
Iter : 2 :P[ 0]=0.2500000 P[ 1]=0.1596875 P[ 2]=0.1437500 P[ 3]=0.1437500
Iter : 3 :P[ 0]=0.1732344 P[ 1]=0.1596875 P[ 2]=0.1437500 P[ 3]=0.1437500
Iter : 4 :P[ 0]=0.1732344 P[ 1]=0.1596875 P[ 2]=0.1111246 P[ 3]=0.1111246
Iter : 5 :P[ 0]=0.1732344 P[ 1]=0.1319559 P[ 2]=0.1111246 P[ 3]=0.1111246
Iter : 6 :P[ 0]=0.1496625 P[ 1]=0.1319559 P[ 2]=0.1111246 P[ 3]=0.1111246
Iter : 7 :P[ 0]=0.1496625 P[ 1]=0.1319559 P[ 2]=0.1011066 P[ 3]=0.1011066
Iter : 8 :P[ 0]=0.1496625 P[ 1]=0.1234406 P[ 2]=0.1011066 P[ 3]=0.1011066
Iter : 9 :P[ 0]=0.1424245 P[ 1]=0.1234406 P[ 2]=0.1011066 P[ 3]=0.1011066
Iter : 10 :P[ 0]=0.1424245 P[ 1]=0.1234406 P[ 2]=0.0980304 P[ 3]=0.0980304
Iter : 11 :P[ 0]=0.1424245 P[ 1]=0.1208259 P[ 2]=0.0980304 P[ 3]=0.0980304
Iter : 12 :P[ 0]=0.1402020 P[ 1]=0.1208259 P[ 2]=0.0980304 P[ 3]=0.0980304
Iter : 13 :P[ 0]=0.1402020 P[ 1]=0.1208259 P[ 2]=0.0970858 P[ 3]=0.0970858
```

```
% ./pgrk 0 -1 verylargegraph.txt
Iter : 4
P[ 0]=0.0136364
P[ 1]=0.0194318
P[ 2]=0.0310227
... other vertices omitted
```

For the HITS algorithm, you need to print two values not one. Follow the convention of the Subject notes

```
Base : 0 :A/H[ 0]=0.3333333/0.3333333 A/H[ 1]=0.3333333/0.3333333 A/H[ 2]=0.3333333/0.3333333
Iter : 1 :A/H[ 0]=0.0000000/0.8320503 A/H[ 1]=0.4472136/0.5547002 A/H[ 2]=0.8944272/0.0000000
```

or for large graphs

```
Iter   : 37
A/H[ 0]=0.0000000/0.0000002
A/H[ 1]=0.0000001/0.0000238
A/H[ 2]=0.0000002/1.0000000
A/H[ 3]=0.0000159/0.0000000
...
```

Deliverables. Include source code of all implemented functions or classes in an archive per Handout 2 guidelines. Document bugs; no bug report no partial points.